


vfio-pci passthrough

Fei Li

fli@suse.com



Summary

- What is VFIO/IOMMU? Why want them?
- VFIO – qemu part
- VFIO – kernel part
- VFIO usage: how to passthrough a pci device





- What is VFIO/IOMMU? Why want them?


What is VFIO/IOMMU?

- The VFIO (Virtual Function I/O) driver is an IOMMU/device agnostic framework for **exposing direct device access to userspace**, in a secure, IOMMU protected environment.
- For x86, it needs the **I/O MMU** hardware support.
- VFIO consists of
 - - kernel device driver: `vfio_pci_driver`, `vfio_iommu_driver`, `vfio_dma`
 - - QEMU device class: `VFIODevice`, `VFIOPCIDevice`
- The guest can operate the pass-throughed PCI device by:
 - accessing the mapped PCI config space and memory space
 - `ioctl()` on a fd of the VFIO kernel device for control operations

What is VFIO/IOMMU? (Continue.)

- In qemu, use VFIO to configure IOMMU:

e.g. `ioctl(VFIO_SET_IOMMU) && ioctl(VFIO_IOMMU_MAP_DMA)`



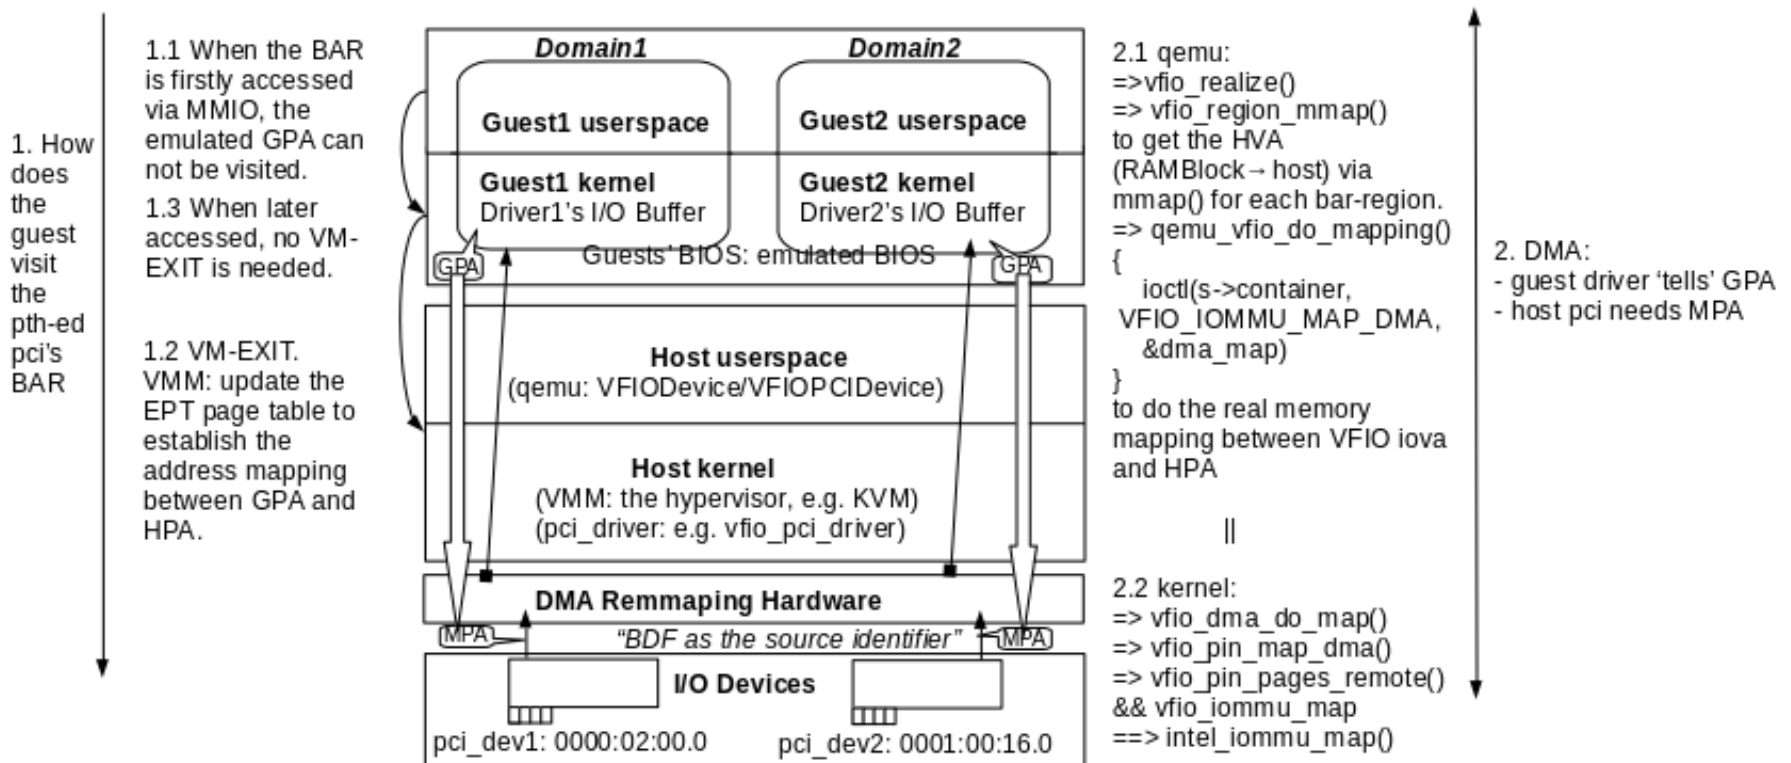
VFIO: device, group, container

- A **group** is a set of devices which is isolatable from all other devices, specialized in **IOMMU**. It is the minimum granularity.
- Within one **container**, different groups can **share a set of page tables** to reduce the duplication. The container provides little functionality: version check and extension query.
- The user needs to add a group into the container.
- The VFIO **device** API includes ioctls for describing the device, the **I/O regions** and their read/write/mmap offsets on the device descriptor, && mechanisms for describing and registering **interrupt** notifications.



Why want them?

- In short, for higher I/O performance by lessening the times of VM-EXIT/VM-ENTRY when accessing PCI BAR and doing DMA.



Why want them? (literal)

- - **when accessing PCI BAR.** The emulated guest BIOS emulate the BAR address for the guest (On the contrary, if expose the host real BAR to the guest, there may be conflicts between host real PCI BAR and the other emulated PCI device's BAR in the guest). For the first time when BAR is accessed and can not be visited, the VM exists and does the address mapping between GPA and HPA using EPT, and records the mapping. When later access, no VM-exit is needed.
- - **when PCI device communicates with GPA via DMA.** When initializing the vfio: `vfio_realize()` in qemu, a memory mapping is established via `vfio_region_mmap()`. Then a `(ioctl(s->container, VFIO_IOMMU_MAP_DMA, &dma_map)` in `qemu_vfio_do_mapping` will do the real mapping between VFIO iova and HPA.



VFIO – qemu part

- How to define vfio-pci in qemu-cli
- Code: how to initialize the vfio-pci device

VFIO – qemu part

1. How to define vfio-pci in qemu command line:

```
`-device vfio-pci, host=0000:02:00.0, bus=pci.1, id=vpci1`
```

2. Code: how to initialize the vfio-pci device

2.1 Deploy the common `device_init_func()` to initialize each qdev in vl.c:

```
main(argc, **argv, **envp) {  
    qemu_opts_foreach(qemu_find_opts("device"),  
                      device_init_func, NULL, &error_fatal);  
}
```

2.2 `device_init_func()` => `qdev_device_add()` to parse the qdev's every attribute (next page)

2.3 `object_property_set_bool()` => ... => `prop->set()` => `property_set_bool` => `device_set_realized()` => `pci_qdev_realize()` => `vfio_realize()`

VFIO – qemu part: qdev_device_add()

```
DeviceState *qdev_device_add(QemuOpts *opts, Error **errp)
{
    BusState *bus = NULL;
    char *driver = qemu_opt_get(opts, "driver");

    /* find driver */
    DeviceClass *dc = qdev_get_device_class(&driver, errp);

    /* find bus */
    char *path = qemu_opt_get(opts, "bus");
    if (path != NULL) {
        bus = qbus_find(path, errp);
    } else if (dc->bus_type != NULL) {
        bus = qbus_find_recursive(sysbus_get_default(), NULL, dc->bus_type);
    }

    /* create device */
    DeviceState *dev = DEVICE(object_new(driver));
    /* add the dev as one of the bus's child && set bus as dev->parent_bus */
    qdev_set_parent_bus(dev, bus);

    qdev_set_id(dev, qemu_opts_id(opts));

    /* set properties */
    if (qemu_opt_foreach(opts, set_property, dev, &err)) {
        goto err_del_dev;
    }

    dev->opts = opts;

    /* call each device/driver's realize(), e.g. vfio_realize() */
    object_property_set_bool(OBJECT(dev), true, "realized", &err);
}
```



VFIO – qemu part: vfio_realize()

```
1. Get the container.  
    /* Create a new container by open "/dev/vfio/vfio" each time. */  
    container = open("/dev/vfio/vfio", O_RDWR);  
  
    if (ioctl(container, VFIO_GET_API_VERSION) != VFIO_API_VERSION)  
        /* Unknown API version handling*/  
    if (!ioctl(container, VFIO_CHECK_EXTENSION, VFIO_X86_IOMMU))  
        /* Doesn't support the IOMMU driver we want. */  
  
2. Get the group, and add the group to the container.  
    /* Get the sysfsdev via `~ device vfio_pci, host=xxxx:xx:xx.x` */  
    vdev->vbasedev.sysfsdev =  
        g_strdup_printf("/sys/bus/pci/devices/%04x:%02x:%02x.%01x",  
            vdev->host.domain, vdev->host.bus, vdev->host.slot, vdev->host.function);  
  
    /* Find the group_id: readlink(the above sysfsdev/iommu_group), e.i.  
    * "/sys/kernel/iommu_groups/14" and basename() to get 14. Then open the group */  
    group = open("/dev/vfio/14", O_RDWR);  
    /* Test the group is viable and available, and some other check */  
    ioctl(group, VFIO_GROUP_GET_STATUS, &group_status);  
  
    /* Add the group to the container */  
    ioctl(group, VFIO_GROUP_SET_CONTAINER, &container);  
  
3. Set the IOMMU via container.  
    /* Enable the IOMMU model we want */  
    ioctl(container, VFIO_SET_IOMMU, VFIO_X86_IOMMU)  
    /* Get addition IOMMU info */  
    ioctl(container, VFIO_IOMMU_GET_INFO, &iommu_info);
```



VFIO – qemu part: vfio_realize() (Continue.)

```
4. Get the pci_dev fd, set the config space and 6 region bars, and the irq info.
/* Get a file descriptor for the device */
device = ioctl(group, VFIO_GROUP_GET_DEVICE_FD, "0000:02:00.0");
/* Test and setup the device */
ioctl(device, VFIO_DEVICE_GET_INFO, &device_info);

for (i = 0; i < device_info.num_regions; i++) {
    struct vfio_region_info reg = { .argsz = sizeof(reg) };
    reg.index = i;
    ioctl(device, VFIO_DEVICE_GET_REGION_INFO, &reg);
    /* Setup mappings... read/write offsets, mmaps. */
    /* For PCI devices, config space is a region */
}
for (i = 0; i < device_info.num_irqs; i++) {
    struct vfio_irq_info irq = { .argsz = sizeof(irq) };
    irq.index = i;
    ioctl(device, VFIO_DEVICE_GET_IRQ_INFO, &reg);
    /* Setup IRQs... eventfds, VFIO_DEVICE_SET_IRQS */
}

5. Map the device iova (e.g. [0, 1MB] as follows) with HPA
/* In qemu, GPA is described by each root MemoryRegion's RAMBlock.
 * For each BAR region, its MR's RAMBlock->host is dma_map.vaddr (HVA) */
dma_map.vaddr = mmap(0, 1024 * 1024, PROT_READ | PROT_WRITE,
                     MAP_PRIVATE | MAP_ANONYMOUS, 0, 0);
dma_map.size = 1024 * 1024;
dma_map.iova = 0; /* 1MB starting at 0x0 from device view */
dma_map.flags = VFIO_DMA_MAP_FLAG_READ | VFIO_DMA_MAP_FLAG_WRITE;
ioctl(container, VFIO_IOMMU_MAP_DMA, &dma_map);

6. /* Gratuitous device reset and go... */
ioctl(device, VFIO_DEVICE_RESET);
```



VFIO – qemu part: qemu_vfio_dma_map()

```
43 struct QEMUVFIOState {
44     QemuMutex lock;
45     /* These fields are protected by BQL */
46     int container, group, device;
47     struct vfio_region_info config_region_info, bar_region_info[6];
48
49     /* These fields are protected by @lock */
50     /* VFIO's IO virtual address space is managed by splitting into a few
51     * sections:
52     *
53     * -----
54     * |xxxxxxxxxxxxxx|
55     * |-----|
56     * | Fixed |
57     * |-----|
58     * |-----|
59     * |-----|
60     * |-----|
61     * | Free |
62     * |-----|
63     * |-----|
64     * |-----|
65     * | Temp |
66     * |-----|
67     * |-----|
68     * |xxxxxxxxxxxxxx|
69     * |xxxxxxxxxxxxxx|
70     * |-----|
71     *
72     * - Addresses lower than QEMU_VFIO_IOVA_MIN are reserved as invalid;
73     *
74     * - Fixed mappings of HVAs are assigned "low" IOVAs in the range of
75     *   [QEMU_VFIO_IOVA_MIN, low_water_mark). Once allocated they will not be
76     *   reclaimed - low_water_mark never shrinks;
77     *
78     * - IOVAs in range [low_water_mark, high_water_mark) are free;
79     *
80     * - IOVAs in range [high_water_mark, QEMU_VFIO_IOVA_MAX) are volatile
81     *   mappings. At each qemu_vfio_dma_reset_temporary() call, the whole area
82     *   is recycled. The caller should make sure I/O's depending on these
83     *   mappings are completed before calling.
84     */
85     uint64_t low_water_mark;
86     uint64_t high_water_mark;
87     IOVAMapping *mappings;
88     int nr_mappings;
89 };
```



VFIO – kernel part

- related drivers
- container/group/device's fops
- connects with pci: `vfio_pci_driver`
- DMA map via `iommu_type1`
- get io: config space + 6 bar
- msix interrupt: `eventfd/irqfd`

VFIO – kernel part: related drivers

1. vfio interface:

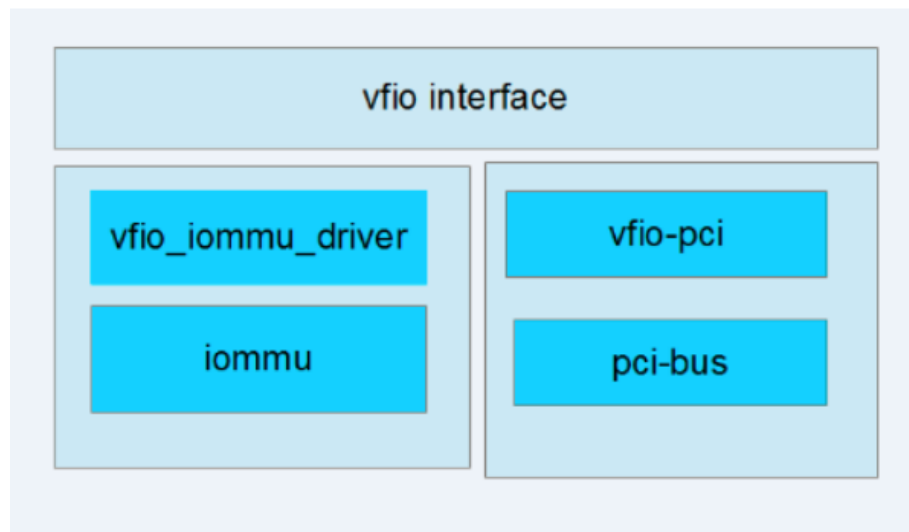
Provide the common ioctl() for userspace, including container, group, device.

2. vfio_iommu_driver (we now use vfio_iommu_type1):

Expose DMA for userspace, e.g. VFIO_IOMMU_MAP_DMA

3. vfio_pci_driver:

Deploy the existed pci standards: to be a pci driver, attach to a pci bus, realize I/O and interrupts, etc.



VFIO – kernel part: container

1. container: via vfio_fops

```
1287 static const struct file_operations vfio_fops = {
1288     .owner          = THIS_MODULE,
1289     .open           = vfio_fops_open,
1290     .release        = vfio_fops_release,
1291     .read           = vfio_fops_read,
1292     .write          = vfio_fops_write,
1293     .unlocked_ioctl = vfio_fops_unl_ioctl,
1294     #ifdef CONFIG_COMPAT
1295     .compat_ioctl   = vfio_fops_compat_ioctl,
1296     #endif
1297     .mmap           = vfio_fops_mmap,
1298 };
```

```
2188 static struct miscdevice vfio_dev = {
2189     .minor = VFIO_MINOR,
2190     .name = "vfio",
2191     .fops = &vfio_fops,
2192     .nodename = "vfio/vfio",
2193     .mode = S_IRUGO | S_IWUGO,
2194 };
2195
2196 static int __init vfio_init(void)
2197 {
2198     int ret;
2199
2200     idr_init(&vfio.group_idr);
2201     mutex_init(&vfio.group_lock);
2202     mutex_init(&vfio.iommu_drivers_lock);
2203     INIT_LIST_HEAD(&vfio.group_list);
2204     INIT_LIST_HEAD(&vfio.iommu_drivers_list);
2205     init_waitqueue_head(&vfio.release_q);
2206
2207     ret = misc_register(&vfio_dev);
```

VFIO – kernel part: group

1. group: via vfio_group_fops

```
1636 static const struct file_operations vfio_group_fops = {
1637     .owner          = THIS_MODULE,
1638     .unlocked_ioctl = vfio_group_fops_unl_ioctl,
1639 #ifdef CONFIG_COMPAT
1640     .compat_ioctl   = vfio_group_fops_compat_ioctl,
1641 #endif
1642     .open           = vfio_group_fops_open,
1643     .release        = vfio_group_fops_release,
1644 };
```

```
2196 static int __init vfio_init(void)
2197 {
2198     ret = misc_register(&vfio_dev);
2199
2200     /* /dev/vfio/$GROUP */
2201     vfio.class = class_create(THIS_MODULE, "vfio");
2202
2203     vfio.class->devnode = vfio_devnode;
2204
2205     ret = alloc_chrdev_region(&vfio.group_devt, 0, MINORMASK, "vfio");
2206
2207     cdev_init(&vfio.group_cdev, &vfio_group_fops);
2208     ret = cdev_add(&vfio.group_cdev, vfio.group_devt, MINORMASK);
```



VFIO – kernel part: device

1. device: indirectly via group fd

To allocate the fd and the struct file which has a anon inode.

```
1558 |     case VFIO_GROUP_GET_DEVICE_FD:
1559 |     {
1560 |         char *buf;
1561 |
1562 |         buf = strndup_user((const char __user *)arg, PAGE_SIZE);
1563 |         if (IS_ERR(buf))
1564 |             return PTR_ERR(buf);
1565 |
1566 |         ret = vfio_group_get_device_fd(group, buf);
1567 |         kfree(buf);
1568 |         break;
1569 |     }
```

```
1480 |         filep = anon_inode_getfile("[vfiio-device]", &vfio_device_fops,
1481 |                                     device, O_RDWR);
```

```
1715 | static const struct file_operations vfio_device_fops = {
1716 |     .owner          = THIS_MODULE,
1717 |     .release        = vfio_device_fops_release,
1718 |     .read           = vfio_device_fops_read,
1719 |     .write          = vfio_device_fops_write,
1720 |     .unlocked_ioctl = vfio_device_fops_unl_ioctl,
1721 | #ifdef CONFIG_COMPAT
1722 |     .compat_ioctl   = vfio_device_fops_compat_ioctl,
1723 | #endif
1724 |     .mmap           = vfio_device_fops_mmap,
1725 | };
```



VFIO – kernel part: vfio_pci_driver

1. `__init vfio_pci_init()` => `pci_register_driver(&vfio_pci_driver)`

=> `bus_add_driver(drv)` => `__driver_attach()` => `really_probe()`

=> `dev->bus->probe()`

[`pci_bus_type.probe = pci_device_probe: pci_drv->probe()`]

=> `vfio_pci_probe`

```
1190 static int vfio_pci_probe(struct pci_dev *pdev,
1191                          const struct pci_device_id *id)
1192 {
1193     group = vfio_iommu_group_get(&pdev->dev);
1194
1195     ret = vfio_add_group_dev(&pdev->dev, &vfio_pci_ops, vdev);
```

```
1315 static struct pci_driver vfio_pci_driver = {
1316     .name           = "vfio-pci",
1317     .id_table       = NULL, /* only dynamic ids */
1318     .probe          = vfio_pci_probe,
1319     .remove         = vfio_pci_remove,
1320     .err_handler    = &vfio_err_handlers,
1321 };
```



VFIO – kernel part: vfio_pci_driver (Continue)

```
805 int vfio_add_group_dev(struct device *dev,
806                       const struct vfio_device_ops *ops, void *device_data)
807 {
808     struct iommu_group *iommu_group;
809     struct vfio_group *group;
810     struct vfio_device *device;
811
812     iommu_group = iommu_group_get(dev);
813
814     group = vfio_group_get_from_iommu(iommu_group);
815     if (!group) {
816         group = vfio_create_group(iommu_group);
817     } else {
818         iommu_group_put(iommu_group);
819     }
820
821     device = vfio_group_get_device(group, dev);
822     if (device) {
823         WARN(1, "Device %s already exists on group %d\n",
824             dev_name(dev), iommu_group_id(iommu_group));
825         ...
826         return -EBUSY;
827     }
828
829     device = vfio_group_create_device(group, dev, ops, device_data);
```

VFIO – kernel part: DMA MAP via iommu_type1

```
1. vfio_fops_unl_ioctl => vfio_iommu_type1_ioctl
=> vfio_dma_do_map
[record (iova, vaddr) in rb_root dma_list of struct vfio_iommu]
=> vfio_pin_map_dma =>
    {1.1 vfio_pin_pages_remote => vaddr_get_pfn => get_user_pages
=> get_user_pages_unlocked => __get_user_pages
=> handle_mm_fault
    1.2 vfio_iommu_map(iova, pfn) => iommu_map()
=> domain→ops→map() [e.g. intel_iommu_map()]
=> domain_pfn_mapping() => __domain_mapping()
=> pfn_to_dma_pte()
    }
```



VFIO – kernel part: get io e.g. config space + 6 bar

```
    } else if (cmd == VFIO_DEVICE_GET_REGION_INFO) {
        struct pci_dev *pdev = vdev->pdev;
        struct vfio_region_info info;
        struct vfio_info_cap caps = { .buf = NULL, .size = 0 };
        int i, ret;

        minsz = offsetofend(struct vfio_region_info, offset);

        if (copy_from_user(&info, (void __user *)arg, minsz))
            return -EFAULT;

        if (info.argsz < minsz)
            return -EINVAL;

        switch (info.index) {
        case VFIO_PCI_CONFIG_REGION_INDEX:
            info.offset = VFIO_PCI_INDEX_TO_OFFSET(info.index);
            info.size = pdev->cfg_size;
            info.flags = VFIO_REGION_INFO_FLAG_READ |
                VFIO_REGION_INFO_FLAG_WRITE;

            break;
        case VFIO_PCI_BAR0_REGION_INDEX ... VFIO_PCI_BAR5_REGION_INDEX:
            info.offset = VFIO_PCI_INDEX_TO_OFFSET(info.index);
            info.size = pci_resource_len(pdev, info.index);
            if (!info.size) {
                info.flags = 0;
                break;
            }

            info.flags = VFIO_REGION_INFO_FLAG_READ |
                VFIO_REGION_INFO_FLAG_WRITE;
            if (vdev->bar_mmap_supported[info.index]) {
                info.flags |= VFIO_REGION_INFO_FLAG_MMAP;
                if (info.index == vdev->msix_bar) {
                    ret = msix_mappable_cap(vdev, &caps);
                    if (ret)
                        return ret;
                }
            }

            break;
        case VFIO_PCI_ROM_REGION_INDEX:
```



VFIO – kernel part: vfio_pci_mmap for each bar

```
1106 static int vfio_pci_mmap(void *device_data, struct vm_area_struct *vma)
1107 {
1108     struct vfio_pci_device *vdev = device_data;
1109     struct pci_dev *pdev = vdev->pdev;
1110     unsigned int index;
1111     u64 phys_len, req_len, pgoff, req_start;
1112     int ret;
1113
1114     index = vma->vm_pgoff >> (VFIO_PCI_OFFSET_SHIFT - PAGE_SHIFT);
1115
1116     if (vma->vm_end < vma->vm_start)
1117         return -EINVAL;
1118     if ((vma->vm_flags & VM_SHARED) == 0)
1119         return -EINVAL;
1120     if (index >= VFIO_PCI_ROM_REGION_INDEX)
1121         return -EINVAL;
1122     if (!vdev->bar_mmap_supported[index])
1123         return -EINVAL;
1124
1125     phys_len = PAGE_ALIGN(pci_resource_len(pdev, index));
1126     req_len = vma->vm_end - vma->vm_start;
1127     pgoff = vma->vm_pgoff &
1128         ((1U << (VFIO_PCI_OFFSET_SHIFT - PAGE_SHIFT)) - 1);
1129     req_start = pgoff << PAGE_SHIFT;
1130
1131     /* Even though we don't make use of the barmap for the mmap,
1132      * we need to request the region and the barmap tracks that. */
1133     if (!vdev->barmap[index]) {
1134         ret = pci_request_selected_regions(pdev,
1135                                           1 << index, "vfio-pci");
1136         if (ret)
1137             return ret;
1138
1139         vdev->barmap[index] = pci_iomap(pdev, index, 0);
1140         if (!vdev->barmap[index]) {
1141             pci_release_selected_regions(pdev, 1 << index);
1142             return -ENOMEM;
1143         }
1144     }
1145
1146     vma->vm_private_data = vdev;
1147     vma->vm_page_prot = pgprot_noncached(vma->vm_page_prot);
1148     vma->vm_pgoff = (pci_resource_start(pdev, index) >> PAGE_SHIFT) + pgoff;
1149 }
```


VFIO – kernel part: interrupt via eventfd/irqfd

If uses MSI-X to do the interrupt:

=> `vfiopci_ioctl()` [parse the struct `vfiopci_irq_set` assigned in qemu
earlier: `vfiopci_msix_vector_do_use()`]

=> `vfiopci_set_irqs_ioctl()` [to check use which irq: intx/msi/msix]

=> `vfiopci_set_msi_trigger()` [if msi or msix]

=> `vfiopci_msi_set_block()` [if uses `VFIO_IRQ_SET_DATA_EVENTFD`]

=> `vfiopci_msi_set_vector_signal()`

[1. //get the eventfd_ctx via the eventfd registered in qemu
eventfd_ctx *trigger = eventfd_ctx_fdget(fd); //fd: next page

2. //allocate an interrupt line: irq for the vfio pci device

```
336         ret = request_irq(irq, vfio_msihandler, 0,  
337                           vdev->ctx[vector].name, trigger);
```

3. //wait for the IRQ occurs, then call `vfiopci_msihandler` to handle.]

```
242 static irqreturn_t vfio_msihandler(int irq, void *arg)  
243 {  
244     struct eventfd_ctx *trigger = arg;  
245     //Notify the userspace which waits on the epoll.  
246     //Ref: qemu_set_fd_handler() to check the qemu part.  
247     eventfd_signal(trigger, 1);  
248     return IRQ_HANDLED;  
249 }
```

(plus) VFIO – qemu part: register the vfio_irq_set

```
int argsz;
struct vfio_irq_set *irq_set;
int32_t *pfd;

argsz = sizeof(*irq_set) + sizeof(*pfd);

irq_set = g_malloc0(argsz);
irq_set->argsz = argsz;
irq_set->flags = VFIO_IRQ_SET_DATA_EVENTFD |
                VFIO_IRQ_SET_ACTION_TRIGGER;
irq_set->index = VFIO_PCI_MSIX_IRQ_INDEX;
irq_set->start = nr;
irq_set->count = 1;
pfd = (int32_t *)&irq_set->data;

if (vector->virq >= 0) {
    *pfd = event_notifier_get_fd(&vector->kvm_interrupt);
} else {
    *pfd = event_notifier_get_fd(&vector->interrupt);
}

ret = ioctl(vdev->vbasedev.fd, VFIO_DEVICE_SET_IRQS, irq_set);
```





VFIO usage: how to passthrough a pci device

VFIO usage: how to passthrough a pci device

1. Enable VT-d in BIOS, and add `intel_iommu=on` in host kernel
2. Build vfio/iommu in kernel or as module
3. Insmod the vfio-pci.ko: `modprobe vfio-pci`
4. If the pci needs to be passthroughed has been bound to other host driver, unbind it firstly:

```
echo 0000:02:00.0 > /sys/bus/pci/devices/0000\:02\:00.0/driver/unbind
```

5. echo “vendor_id device_id” > /sys/bus/pci/drivers/vfio-pci/new_id

PS: `lspci -n -s BDF` to check the vendor_id and device_id

```
linux-50ts:/build/gitcode # lspci -n -s 00:02.0  
00:02.0 0300: 8086:191b (rev 06)
```

Reference

- <https://www.kernel.org/doc/Documentation/vfio.txt>
- <https://lwn.net/Articles/509153/>
- kernel source code
- qemu source code



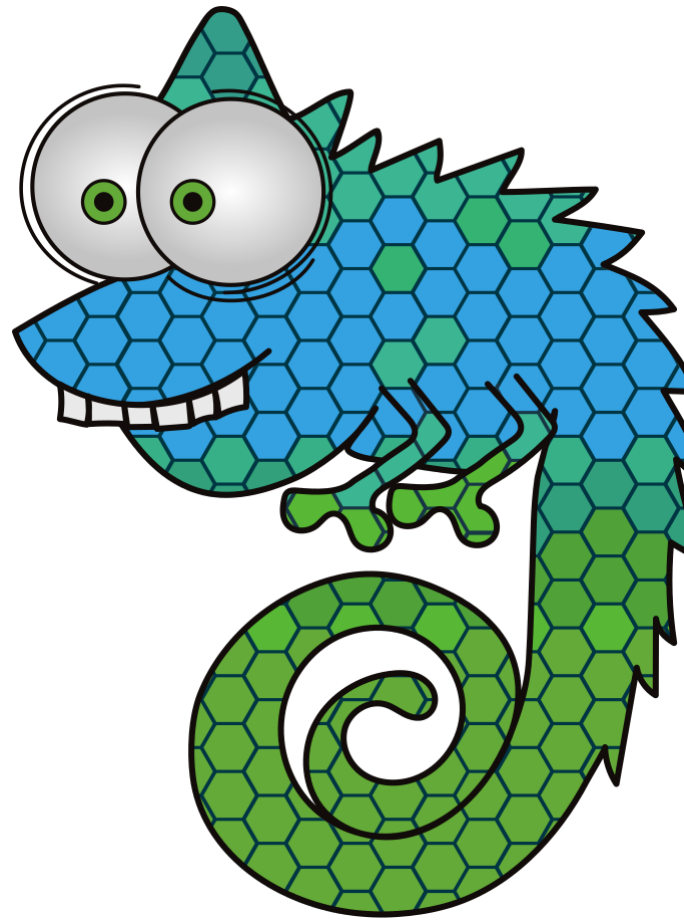


Questions?

Join the conversation,
contribute & have a lot of fun!
www.opensuse.org

Thank you.





Have a Lot of Fun, and Join Us At:

www.opensuse.org



License

This slide deck is licensed under the Creative Commons Attribution-ShareAlike 4.0 International license. It can be shared and adapted for any purpose (even commercially) as long as Attribution is given and any derivative work is distributed under the same license.

Details can be found at <https://creativecommons.org/licenses/by-sa/4.0/>

General Disclaimer

This document is not to be construed as a promise by any participating organisation to develop, deliver, or market a product. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. openSUSE makes no representations or warranties with respect to the contents of this document, and specifically disclaims any express or implied warranties of merchantability or fitness for any particular purpose. The development, release, and timing of features or functionality described for openSUSE products remains at the sole discretion of openSUSE. Further, openSUSE reserves the right to revise this document and to make changes to its content, at any time, without obligation to notify any person or entity of such revisions or changes. All openSUSE marks referenced in this presentation are trademarks or registered trademarks of SUSE LLC, in the United States and other countries. All third-party trademarks are the property of their respective owners.

Credits

Template

Richard Brown
rbrown@opensuse.org

Design & Inspiration

openSUSE Design Team
<http://opensuse.github.io/branding-guidelines/>